Alec Stashevsky

alecstashevsky.com

alec@alecstashevsky.com

Applied machine learning scientist and technical leader with a track record of leading teams shipping machine learning / AI solutions at scale. Leader of Fetch's core machine learning team building solutions powering the Fetch app handling over \$150 billion in annual gross merchandise volume, an equivalent to the third largest retailer in the United States. Technical expertise in deep learning for document AI, semantic search/entity resolution, ads ranking, recommendations systems, and fraud detection.

PROFESSIONAL EXPERIENCE

Head of Relevance and Discovery / Fetch

(10/2024)-Present

- Leading four teams responsible for Core Ranking, Ranking Systems, ML Platform Engineering, and Measurement.
- Engineering manager and technical leader for 30+ engineers in Discovery group. Building teams and ML-driven search + discovery features from 0 to 1.
- Building core ML infrastructure to support large scale real-time workloads including streaming data aggregation via Apache Kafka and Flink, feature store, GPU cluster management and deployment, and hybrid search infrastructure.

Tech Lead Manager, Core Machine Learning / Fetch (10/2022)-(10/2024)

- Lead a team of 15+ scientists and engineers building world-class ML/AI technology powering the core of the Fetch app. Our systems extract information from over 11 million receipts in real-time every day, and process over \$150 billion in gross merchandise volume annually.
- Core contributor and technical leader for Fetch's largest product launch to date with over \$200M in annual revenue attributable to our core document Al technology.
- Research, build, and deploy deep neural networks touching computer vision, natural language processing, and graph machine learning.
- Pre-training and fine-tuning of large language models (LLMs), vision encoder-decoders, and graph neural networks.
- Engineering leader supporting AI/ML microservices, edge device deployments, and serving optimizations, including quantization, compilation, pruning, and knowledge distillation.
- Lead partnerships with open-source and academic communities including Stanford University,
 Hugging Face, PyTorch, PyTorch Geometric, AWS SageMaker, and Streamlit.

Senior Machine Learning Scientist / Fetch

(10/2021)-(10-2022)

- Research and develop document understanding models including CNNs and Vision Transformers for optical character recognition, large language transformers, and graph neural networks for named entity recognition and key information extraction.
- Use tools like PyTorch, Hugging Face, PyTorch Geometric, scikit-learn, SQL, R, Docker, AWS and Snowflake.

Financial Economist / NERA Economic Consulting

(09/2020)-(09/2021)

- Look inside the books of some of the largest financial institutions in the world to estimate damages and predict the performance of complex financial instruments responsible for highprofile banking crises, securities fraud, and market-manipulation cases.
- Build probabilistic financial models using Markov chain Monte Carlo methods, random matrix theory, and time-series forecasting.
- Identify, explain, and value litigation involving mortgage-backed securities (RMBS), collateralized debt obligations (CDOs), swaps, and other derivatives underpinning trillions of dollars in assets.
- Provide evidence and expert economic testimony for Fortune 500 companies, SEC, DOJ, and FINRA.

Economist / The Cadmus Group

(11/2019) - (09/2020)

- Lead the design and evaluation of demand-side management programs, including a \$600k+ randomized control trial on smart thermostat direct load-control.
- Forecast electric vehicle adoption, demand elasticity, and electrification for budgeting hundreds of millions of dollars under diverse energy industry clients' management.
- Design difference-in-difference models, demand-elasticity programs, and causal inference mechanisms to provide gold-standard reporting to regulators and operators responsible for most of the United States energy supply.

Data Scientist/Kyrgies

(08/2017) - (10/2019)

- Designed marketing mix models, audience creation, and targeted advertising solutions for a bootstrapped e-commerce startup.
- Deployed recommendation systems and probabilistic engagement models (pCTR, pConversion)
 for targeted advertising that led to 8% reduction in CAC.

PUBLICATIONS AND TALKS

- Speaker, "Optimizing Sentence Transformers for Entity Resolution at Scale." *Convergence*, May 2024. (recording)
- Speaker, "Graphing Groceries: Understanding Receipts with Transformer Hydranets." *Microsoft AI Graph Learning Group*, January 2024. (recording)
- Speaker, "Graph Transformers for Semantic Link Prediction at Scale." *Stanford Graph Learning Workshop*, October 2023. (recording)
- Assessing Strategies to Reduce the Carbon Footprint of the Annual Meeting of the American Academy of Ophthalmology. *JAMA Ophthalmology*, August 2023. doi:10.1001/jamaophthalmol.2023.3516
- Speaker, "Production Graph ML at Fetch." PyTorch Geometric Town Hall, January 2023. (recording)
- Teaching to Our Time: a Survey Study of Current Opinions and Didactics About Climate Mental Health Training in US Psychiatry Residency and Fellowship Programs. *Academic Psychiatry*, July 2022. doi:10.1007/s40596-022-01680-7
- Estimation of the Carbon Footprint Associated With Attendees of the American Psychiatric Association Annual Meeting. *JAMA Network Open*, January 2021. doi:10.1001/jamanetworkopen.2020.35641

ORGANIZATIONS

Statistician / Group for the Advancement of Psychiatry, Climate Committee (06/2020)-Present

GAP is a think tank of top psychiatric minds whose thoughtful analysis and recommendations serve to
influence and advance modern psychiatric theory and practice. I serve as a resident statistician
supporting climate-related research (see publications above).

OPEN-SOURCE SOFTWARE DEVELOPMENT

• Author and maintainer of {blocklength}, an R package for advanced statistical analyses of time-series and other dependent data with over 12 thousand downloads. {blocklength} provides tooling to optimize the block-length parameter of a dependent bootstrap procedure (the block-bootstrap) commonly used in time-series modelling. Available on CRAN and at my GitHub.

EDUCATION

B.A. Mathematics-Economics / Reed College, Portland OR